

Randomized algorithms

Randomness in algorithms


- ▶ Data is random \Rightarrow average-case analysis of algorithms
- ▶ Algorithm is non-deterministic, uses randomness \Rightarrow **randomized algorithms**

Some applications of randomization

- ▶ **Universal hashing**: helps to avoid dependency on data distribution, insures security (e.g. adversary-made worst cases)
- ▶ **Locality-sensitive hashing**: helps “sketching”, i.e. representing big objects by small objects without much distortion of distances
- ▶ **Online algorithms**: improves the competitiveness
- ▶ **Approximate counting and other streaming algorithms**: drastically reduces spaces


Accessing an exclusive resource by non-communicating processes

- ▶ n non-communicating processes want to access a resource in rounds
- ▶ if two processes ask for access at the same round, none of them gets access
- ▶ *What would be the good strategy?*



Accessing an exclusive resource by non-communicating processes

- ▶ n non-communicating processes want to access a resource in rounds
- ▶ if two processes ask for access at the same round, none of them gets access
- ▶ *What would be the good strategy?*
- ▶ Each process ask access with probability p (“*symmetry breaking*”)
- ▶ optimum: $p = 1/n$
- ▶ then each process will get access on average between $1/en$ and $1/2n$ of time (depending on n)
- ▶ with probability at least $(1 - 1/n)$ all processes access R at least once in $O(n \log n)$ rounds (about $2en \cdot \ln(n)$)



Besides ...



- ▶ ... randomized algorithms are usually *simple* (to implement, not to understand) and, as a consequence, *efficient*

Two types of randomized algorithms

- ▶ **Las Vegas**: always correct (but fast *in expectation*, e.g. *Randomized QUICKSORT*)
- ▶ **Monte Carlo**: admits a small error probability

Books on randomized algorithms

- ▶ Motwani, Raghavan, Randomized algorithms, Cambridge University Press, 1995
- ▶ Mitzenmacher, Upfal, Probability and Computing: Randomized Algorithms and Probabilistic Analysis, Cambridge University Press, 2005



Monte Carlo: simple example

- ▶ Given a (very large) array of numbers, return some array element larger than the median

Edge connectivity (global min-cut)


- ▶ Given an undirected graph $G = (V, E)$, a *cut* is $V = A \cup B$ ($A \cap B = \emptyset$). What the minimal number of edges connecting the two parts of a cut of G ?
- ▶ Can be solved in polynomial time:
 - ▶ $O(n^2m)$ using Edmons-Karp
 - ▶ $O(n(n + m) \log n)$ using Stoer-Wagner
- ▶ A much simpler *randomized* algorithm exists due to David Karger (1992)

High-level idea

- ▶ Stoer-Wagner takes time $O((n + m) \log n)$ to find the pair of nodes to merge
- ▶ *Idea*: pick a random edge and merge endpoints (contract the edge)

Contraction algorithm

- ▶ Like Stoer-Wagner, the algorithm works with *multigraphs*.
(The input graph G can be a multi-graph as well)




$$C=\{c,d,e,j,k,l\}=(\{1,5\},\{2,3,4\})$$

$$C=\{e,f,g\}=(\{2\},\{1,3,4,5\})$$

Contraction

- ▶ Contraction of an edge (u, v) :
 - ▶ merge u and v into a single new node w
 - ▶ for all edges (t, u) or (t, v) , update to (t, w)
 - ▶ keep parallel edges, but delete self-loops



- ▶ *Remark (cf Stoer-Wagner):* Contraction can not decrease the min-cut value (but can increase)

Contraction algorithm

do


pick an edge $e = (u, v)$ uniformly at random;

contract (u, v)


until G has just two nodes s, t

return the corresponding cut (nodes contracted to s and t respectively) and the number of crossing edges


Example




Example




Example




Example




Example



Example



Example




Analysis of the algorithm

Theorem: The contraction algorithm returns a min cut with probability $\geq \frac{2}{n(n-1)}$

Proof: Consider a global min-cut (A^*, B^*) of G . Let F^* be edges with one endpoint in A^* and the other in B^* . Let $k = |F^*|$ = size of min cut.

- ▶ In first step, algorithm contracts an edge in F^* with probability $k/|E|$.
- ▶ Every node has degree $\geq k$ since otherwise (A^*, B^*) would not be min-cut $\Rightarrow |E| \geq \frac{kn}{2}$.
- ▶ Thus, algorithm contracts an edge in F^* with probability $\leq \frac{2}{n}$



Analysis of the algorithm (cont)

- ▶ Let G' be graph after j iterations. There are $n' = n - j$ supernodes.
- ▶ Suppose no edge in F^* has been contracted. The min-cut in G' is still k .
- ▶ Since value of min-cut is k , $|E'| \geq kn'/2$.
- ▶ Thus, algorithm contracts an edge in F^* with probability $\leq 2/n'$.
- ▶ Let E_j = event that an edge in F^* is not contracted in iteration j .

$$\begin{aligned} P[E_1 \cap E_2 \cap \dots E_{n-2}] &= P[E_1] \cdot P[E_2|E_1] \cdot \dots \cdot P[E_{n-2}|E_1 \cap E_2 \cap \dots E_{n-3}] \\ &\geq \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \dots \left(1 - \frac{2}{4}\right) \left(1 - \frac{2}{3}\right) \\ &= \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \dots \left(\frac{2}{4}\right) \left(\frac{1}{3}\right) \\ &= \frac{2}{n(n-1)} \geq \frac{2}{n^2} \end{aligned}$$

Contraction algorithm: amplification

Amplification: To amplify the probability of success, run the contraction algorithm many times and take the best cut!

Claim: If we repeat the contraction algorithm $n^2 \ln n$ times with independent random choices, the probability of failing to find the global min-cut is at most $1/n^2$.

Proof: By independence, the probability of failure is at most

$$\left(1 - \frac{2}{n^2}\right)^{n^2 \ln n} = \left(\left(1 - \frac{2}{n^2}\right)^{\frac{n^2}{2}}\right)^{2 \ln n} \leq (e^{-1})^{2 \ln n} = \frac{1}{n^2}$$

↑
 $(1 - 1/x)^x \leq 1/e$

Contraction algorithm: complexity

- ▶ Contraction algorithm: $O(m \cdot n^2 \cdot \log(n)) = O(n^4 \cdot \log(n))$
- ▶ Solution by network flow requires time $O(n^4)$
- ▶ However, the contraction algorithm can be improved to $O(m \cdot \log^3(n))$ [Karger, Stein, Journal of the ACM, 1996; Karger, Journal of the ACM, 2000]