



Hashing



Plan

- ▶ "Classical" hashing
 - ▶ hashing by chaining
 - ▶ hashing by open addressing
- ▶ Universal hashing
- ▶ Perfect hashing (quick review)
- ▶ Cuckoo hashing
- ▶ Bloom filters
- ▶ Locality-sensitive hashing

Example 1: path finding

- ▶ Assume you want to implement A* shortest path search on a graph of 1000 nodes. You can allocate an array of size 1000 to store distances
- ▶ What about search on Rubik's cube graph (order of 10^6 for $2 \times 2 \times 2$ cube, 10^{19} for $3 \times 3 \times 3$ cube)?

Example 2: data bases

- ▶ Maintain a set of employees (students, messenger users, ...), each identified by a social security number (student ID, phone number, ...)

Example 3: deduplication

- ▶ In a programming language compiler, how to store user-declared identifiers?
- ▶ Construct an index of a book with all terms pointing to their first occurrence in the book

Hash tables: supported operations

- ▶ A generalization of arrays (“direct addressing”)
- ▶ *Goal*: maintain a (possibly evolving) set of objects belonging to a large “universe” (e.g. configurations, ID numbers, words, etc.)
- ▶ *Applications*: deduplication, indexing, path finding, file integrity test (checksum), etc.

Hash tables: supported operations

- ▶ A generalization of arrays (“direct addressing”)
 - ▶ *Goal*: maintain a (possibly evolving) set of objects belonging to a large “universe” (e.g. configurations, ID numbers, words, etc.)
 - ▶ *Applications*: deduplication, indexing, path finding, file integrity test (checksum), etc.
 - ▶ *INSERT*: add a new object
 - ▶ *DELETE*: delete existing object
 - ▶ *LOOKUP*: check for an object
- } possibly specified by a *key*
"associative array"

“Dictionary” data structure

Naive solutions

▶ Bit array (bitmap)

- ▶ still too big for huge applications
- ▶ does not support access to objects
- ▶ BUT ... (cf Bloom filters at the end of this lecture)

▶ Linked list

- ▶ look-up too slow

▶ Search trees

- ▶ better but still slow and memory demanding


Hash tables

► Notation

- U : universe of all possible keys (Ex: strings, IP addresses, game configurations, ...)
- K : subset of keys (actually stored in the dictionary), $|K| \ll |U|$
- $|K| = n$

► Use a table of size proportional to $|K|$: **hash table**

- we lose the direct-addressing ability
- **hash function** maps keys to entries of the hash table (**buckets** or slots)



Hash functions




- ▶ **Hash function h :** Mapping from U to the slots of a hash table $T[0..m-1]$.

$$h : U \rightarrow \{0, 1, \dots, m-1\}$$

- ▶ With direct addressing, key k maps to slot $A[k]$
- ▶ With hash tables, key k maps or “**hashes**” to bucket $T[h[k]]$
- ▶ $h[k]$ is the **hash value** (or simply **hash**) of key k

Hashing and collisions




Collisions: birthday "paradox"

- ▶ What is the probability that two people from this class (100 students) have their birthday the same day?


Collisions: birthday "paradox"

- ▶ What is the probability that two people from this class (100 students) have their birthday the same day?
- ▶ Answer: ≈ 0.9999997
- ▶ Birthday paradox: in a group of 23 people, there is about 50% chance that two people have the same birthday
- ▶ 40 people: 89%, 60 people: 99.4%
- ▶ Conclusion: collisions are frequent

I. Collision Resolution by Chaining



Collision Resolution by Chaining



Hashing with chaining

- ▶ $\text{INSERT}(T.k) : O(1)$
- ▶ $\text{DELETE}(T.k), \text{LOOKUP}(k) : O(\text{list length})$
- ▶ \Rightarrow a good hash function should distribute keys into buckets *as uniformly as possible*
- ▶ random hashing \Rightarrow expected list length is $\alpha = n/m$ (load factor)
- ▶ the *average time* of DELETE and LOOKUP is $O(1 + \alpha)$
 $\Rightarrow O(1)$ if $n = O(m)$ (practical case)

Good hash functions

- ▶ Hash function should be easy to compute
- ▶ Designing good hash functions is tricky. It is easy to design a bad hash function
- ▶ *Examples*: Phone numbers. Benford's law (e.g. prices, population sizes, ...)
- ▶ Keys are usually considered as natural numbers
- ▶ *Example*: Interpret a character string as an integer expressed in some radix notation. Suppose the string is CLRS:
 - ▶ ASCII values: C=67, L=76, R=82, S=83.
 - ▶ There are 128 basic ASCII values.
 - ▶ So, $CLRS = 67 \cdot 128^3 + 76 \cdot 128^2 + 82 \cdot 128^1 + 83 \cdot 128^0 = 141,764,947$

Division Method

- ▶ Map a key k into one of the m slots by taking the remainder of k divided by m . That is,

$$h(k) = k \bmod m$$

- ▶ *Example:* $m = 31$ and $k = 78 \Rightarrow h(k) = 16$
- ▶ *Advantage:* Fast, since requires just one division operation
- ▶ *Disadvantage:* Have to avoid certain values of m .
 - ▶ Don't pick certain values, such as $m = 2^p$ (as the hash won't depend on all bits of k)
- ▶ *Good choice for m :*
 - ▶ Primes, not too close to power of 2 (or 10) are good.

Multiplication Method


- ▶ If $0 < A < 1$, $h(k) = \lfloor m (kA \bmod 1) \rfloor = \lfloor m(kA - \lfloor kA \rfloor) \rfloor$ where $(kA \bmod 1) = kA - \lfloor kA \rfloor$: the fractional part of kA
- ▶ *Disadvantage*: Slower than the division method.
- ▶ *Advantage*: Value of m is not critical.
 - ▶ Typically chosen as a power of 2, i.e., $m = 2^p$, which makes the implementation easy
- ▶ *Example*: $m = 1000, k = 123, A = 0.6180339887\dots$
$$h(k) = \lfloor 1000(123 \cdot 0.6180339887 \bmod 1) \rfloor = \lfloor 1000 \cdot 0.018169\dots \rfloor = 18$$

II. Collision Resolution by Open Addressing

- ▶ All elements are stored in the hash table itself
- ▶ $\Rightarrow n \leq m$, no pointers
- ▶ hash function $h(k, i)$ where $i = 0, 1, 2, \dots, m - 1$, and $\langle h(k, 0), h(k, 1), \dots, h(k, m - 1) \rangle$ is a permutation
- ▶ when inserting/looking up k , probe $h(k, 0), h(k, 1), \dots$ (probe sequence) until
 - ▶ we find k , or
 - ▶ the bucket contains *nil*, or
 - ▶ m buckets have been unsuccessfully probed
- ▶ deletion is complicated, needs a special key "deleted", time may not be dependent on the load factor

Open Addressing: Linear probing


- ▶ The colliding item is placed in a different cell of the table
 - ▶ **Linear probing** handles collisions by placing the colliding item in the next (circularly) available table cell
 $h(k, i) = (h'(k) + i) \bmod m$
 - ▶ Each table cell inspected is referred to as a “probe”
 - ▶ Colliding items clump together, causing future collisions to cause a longer sequence of probes
- ▶ *Example:*
- ▶ $h'(k) = k \bmod 13$
 - ▶ Insert keys
18, 41, 22, 44, 59, 32, 31, 73,
in this order



Example (cont.)

$$h'(k) = k \bmod 13$$

18 41 22 44 59 32 31 73



Quadratic probing

- ▶ $h(k, j) = (h'(k) + c_1 \cdot j + c_2 \cdot j^2) \bmod m$
- ▶ for example, $h(k, j) = (h'(k) + \frac{1}{2} \cdot j + \frac{1}{2} \cdot j^2) \bmod m$
 $h(k, 0), h(k, 1), \dots, h(k, m - 1)$ is a permutation if m is a power of 2
- ▶ quadratic probing works better than linear probing (less clumping)

Double Hashing

- ▶ Double hashing uses a secondary hash function $d(k)$ and handles collisions by placing an item in the first available cell of the series

$$h(k, j) = (h(k) + jd(k)) \bmod m$$


for $j = 0, 1, \dots, m-1$

- ▶ The secondary hash function $d(k)$ cannot have zero values
- ▶ m should be relatively prime to $d(k)$, e.g. $m = 2^q$ and $d(k)$ odd, or m is prime and $d(k) < m$
- ▶ Double hashing is usually more efficient than linear and quadratic probing

Example of Double Hashing

- ▶ Consider a hash table storing integer keys that handles collision with double hashing
 - ▶ $m = 13$
 - ▶ $h(k) = k \bmod 13$
 - ▶ $d(k) = 7 - k \bmod 7$
- ▶ Insert keys 18, 41, 22, 44, 59, 32, 31, 73, in this order

k	$h(k)$	$d(k)$	Probes	
18	5	3	5	
41	2	1	2	
22	9	6	9	
44	5	5	5	10
59	7	4	7	
32	6	3	6	
31	5	4	5	9 0
73	8	4	8	



Performance of Open Addressing

- ▶ Assuming that $\langle h(k, 0), h(k, 1), \dots, h(k, m - 1) \rangle$ is a random permutation (uniformly drawn), the expected number of probes in an insertion (or unsuccessful search) with open addressing is

$$1/(1 - \alpha),$$

where $\alpha = n/m$ the load factor

- ▶ *Explanation:*

let $p_i = P[i \text{ first buckets are full}] = \alpha^i \quad (p_0 = 0)$

$E[\text{number of probes}] = 1 +$

$\sum_{i=1..m-1} i \cdot P[i \text{ full buckets followed by an empty one}] = 1 +$

$\sum_{i=1..m-1} i \cdot (p_{i-1} - p_i) = 1 + \sum_{i=1..m-1} p_i \approx$

$$1 + \alpha + \alpha^2 + \alpha^3 + \dots = 1/(1 - \alpha)$$

Performance of Open Addressing


- ▶ Assuming that $\langle h(k, 0), h(k, 1), \dots, h(k, m - 1) \rangle$ is a random permutation (uniformly drawn), the expected number of probes in an insertion (or unsuccessful search) with open addressing is

$$1/(1 - \alpha),$$

where $\alpha = n/m$ the load factor

- ▶ The expected number of probes for a successful search is
$$(1/\alpha) \log(1/(1 - \alpha))$$

Historical remarks




Hashing by open addressing:
Analysed by Donald Knuth in 1962 (invention attributed
to Andrei Ershov)

Hashing: some conclusions


- ▶ Chaining:
 - ▶ easy implementation
 - ▶ fast in practice
 - ▶ uses more memory
- ▶ Open addressing:
 - ▶ uses less memory
 - ▶ more complex removals
- ▶ Implemented in standard libraries, e.g. `std::unordered_map` in C++

Linear probing vs Double hashing

comparison of average number of operations



successful search



unsuccessful search

Exercise

- ▶ Consider inserting the keys 10, 22, 31, 4, 15, 28, 17, 88, 59 into a hash table of length $m = 11$ using open addressing with the primary hash function $h'(k) = k \bmod m$. Illustrate the result of inserting these keys using linear probing, using quadratic probing with $c_1 = 1$ and $c_2 = 3$, and using double hashing with $h_2(k) = 1 + (k \bmod (m - 1))$.